

Technical Report: MNIST Diffusion Experiments on Purdue Gautschi

PurdueHPC_Codex Workflow

February 24, 2026

1 Scope and Environment

This report summarizes end-to-end execution of diffusion model experiments on Purdue RCAC Gautschi, including environment setup, Slurm execution, model details, and resulting artifacts. All commands and scripts were executed under:

- Scratch workspace: `/scratch/gautschi/rmaulik/codex_test`
- Source repository: `PurdueHPC_Codex`
- Account: `rmaulik`
- Partition: `ai`

No credentials or authentication secrets are stored in this report.

2 Allocation and Job Accounting

Current allocation snapshot from `slist` on Gautschi:

- AI partition GPU-hour balance: **43,798.5**
- CPU partition balance: **0**

Completed experiment jobs:

Job ID	Name	State	Elapsed (s)	GPUs
8109439	mnist_ddpm	COMPLETED	41	1
8109477	mnist_ddpm_long	COMPLETED	99	1
8109577	mnist_ddpm_long	COMPLETED	411	1

Estimated direct GPU-hours from these three runs:

$$\text{GPU-hours} = \frac{41 + 99 + 411}{3600} \times 1 = 0.153 \text{ GPU-hours (approx.)}$$

3 Diffusion Model Formulation

The training script implements a DDPM-style denoising objective over MNIST ($x_0 \in [-1, 1]^{1 \times 28 \times 28}$).

3.1 Forward Diffusion

Define variance schedule $\{\beta_t\}_{t=1}^T$ with

$$\beta_t \in [10^{-4}, 2 \times 10^{-2}], \quad \alpha_t = 1 - \beta_t, \quad \bar{\alpha}_t = \prod_{s=1}^t \alpha_s.$$

The forward Markov transition is

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t} x_{t-1}, \beta_t I),$$

with closed form sample

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I).$$

3.2 Reverse Model and Objective

The denoiser network $\epsilon_\theta(x_t, t)$ predicts injected noise. Reverse mean step used in sampling:

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right).$$

At training time, the objective is standard noise-prediction MSE:

$$\mathcal{L}(\theta) = \mathbb{E}_{x_0, \epsilon, t} \left[\left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|_2^2 \right].$$

4 Network Architecture

The current model is a residual U-Net style denoiser with sinusoidal time embeddings:

- Stem: Conv 1 → 64
- Down path: residual blocks at 64 and 128 channels with strided downsampling (28 → 14 → 7)
- Bottleneck: two residual blocks at 256 channels
- Up path: transposed convolutions + skip concatenation + residual blocks (7 → 14 → 28)
- Output: GroupNorm + Conv 64 → 1
- Optimizer: AdamW, gradient clipping (max norm 1.0)

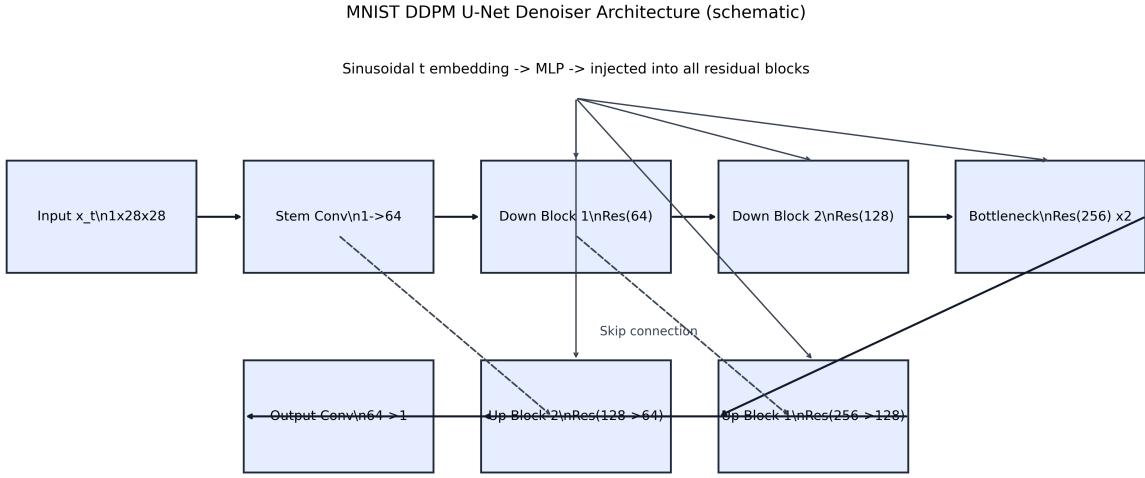


Figure 1: U-Net denoiser schematic generated by the training pipeline.

5 Latest Long Run Configuration

The latest long run (job 8109577) used:

- Epochs: 80
- Diffusion steps (T): 300
- Batch size: 128
- Slurm resources: 1 H100 GPU, 14 CPUs, partition `ai`, walltime 4h

From `metrics.json` (run `long_8109577`):

- Total steps: 37,520
- Final epoch mean loss: 0.0362516
- Best epoch mean loss: 0.0360608
- Device: CUDA

6 Results and Artifacts

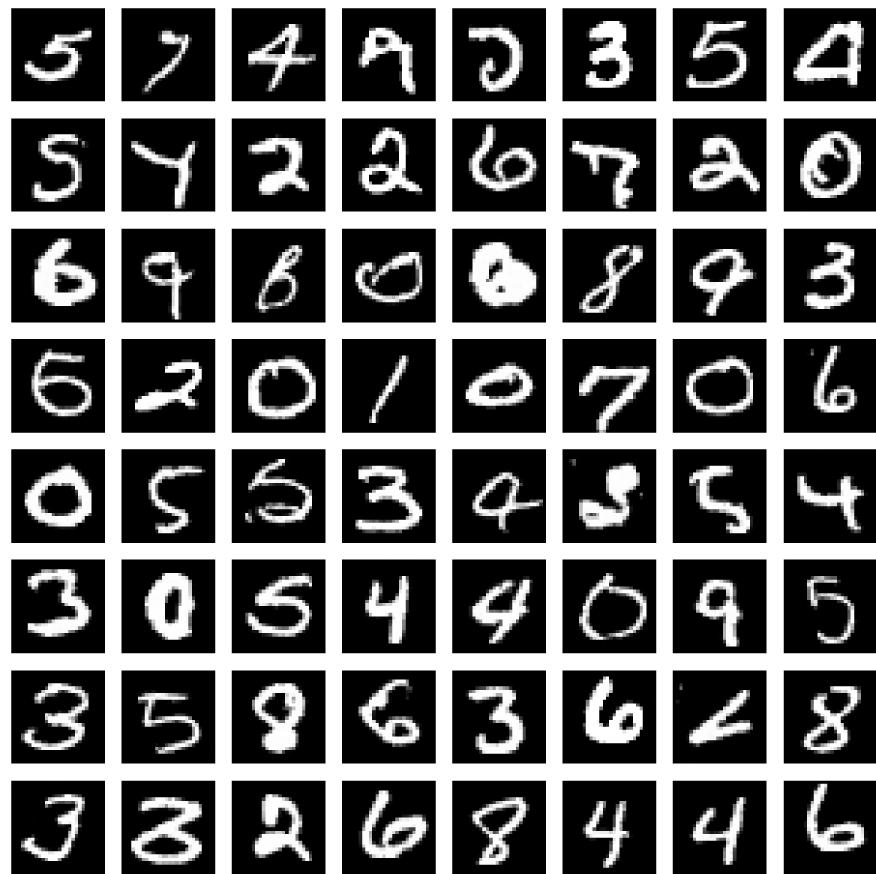


Figure 2: Generated MNIST samples from the latest completed run.

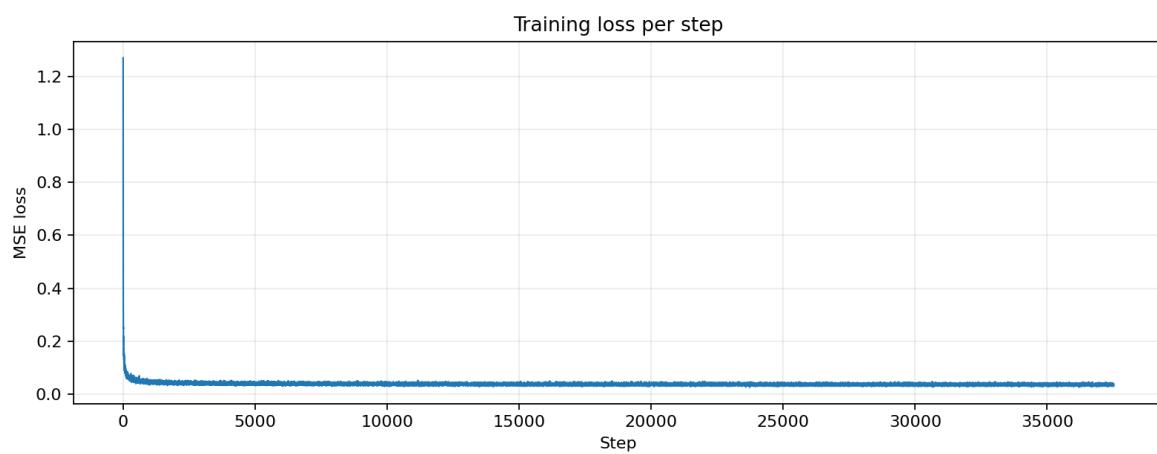


Figure 3: Per-step training loss trajectory.




Figure 4: Per-epoch mean loss trajectory.

7 Posterior Sampling with Partial Observations

We added a post-training posterior sampling routine to `mnist_diffusion.py` that performs likelihood-guided reverse diffusion under partial pixel observations. For this experiment:

- Observed fraction: 70% of pixels (30% occluded, random mask)
- Target digit class: 7
- Guidance scale: 1.5
- Likelihood noise scale: 0.1

Validated Gautschi posterior job:

- Job ID: 8112300
- State: COMPLETED, exit code 0:0
- Elapsed: 00:00:28
- Run tag: `posterior_70pct_8112300`




Figure 5: Posterior conditioning setup: ground truth, observed mask (70% observed), observed pixels, and posterior draws.



Figure 6: Likelihood-guided posterior samples for the partially observed digit.

8 Live Dashboard Operation

The workflow now maintains a root live dashboard at:

- `outputs/dashboard.html`

This dashboard follows the latest run via:

- `outputs/LATEST_RUN.txt`
- `outputs/current` symlink

and supports:

- periodic refresh during active runs,
- automatic stop of refresh when status reaches `completed`,
- image zoom controls for plots/schematics.

9 Reproducibility Notes

Primary scripts tracked in Git:

- `mnist_diffusion.py`
- `submit_mnist_diffusion_long.slurm`
- `serve_dashboard.sh`

To regenerate this report PDF locally:

```
cd PurdueHPC_Codex/report
latexmk -pdf -interaction=nonstopmode -halt-on-error \
-output-directory=../output/pdf gautschi_diffusion_report.tex
```