

# Technical Report: MNIST Diffusion Experiments on Purdue Gautschi

PurdueHPC\_Codex Workflow

March 1, 2026

## 1 Scope and Environment

This report summarizes end-to-end execution of diffusion model experiments on Purdue RCAC Gautschi, including environment setup, Slurm execution, model details, and resulting artifacts. All commands and scripts were executed under:

- Scratch workspace: `/scratch/gautschi/rmaulik/codex_test`
- Source repository: `PurdueHPC_Codex`
- Account: `rmaulik`
- Partition: `ai`

No credentials or authentication secrets are stored in this report.

## 2 Allocation and Job Accounting

Current allocation snapshot from `slist` on Gautschi:

- AI partition GPU-hour balance: **43,798.5**
- CPU partition balance: **0**

Completed experiment jobs:

Job ID	Name	State	Elapsed (s)	GPUs
8109439	mnist_ddpm	COMPLETED	41	1
8109477	mnist_ddpm_long	COMPLETED	99	1
8109577	mnist_ddpm_long	COMPLETED	411	1

Estimated direct GPU-hours from these three runs:

$$\text{GPU-hours} = \frac{41 + 99 + 411}{3600} \times 1 = 0.153 \text{ GPU-hours (approx.)}$$

## 3 Diffusion Model Formulation

The training script implements a DDPM-style denoising objective over MNIST ( $x_0 \in [-1, 1]^{1 \times 28 \times 28}$ ).

### 3.1 Forward Diffusion

Define variance schedule  $\{\beta_t\}_{t=1}^T$  with

$$\beta_t \in [10^{-4}, 2 \times 10^{-2}], \quad \alpha_t = 1 - \beta_t, \quad \bar{\alpha}_t = \prod_{s=1}^t \alpha_s.$$

The forward Markov transition is

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t} x_{t-1}, \beta_t I),$$

with closed form sample

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I).$$

### 3.2 Reverse Model and Objective

The denoiser network  $\epsilon_\theta(x_t, t)$  predicts injected noise. Reverse mean step used in sampling:

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right).$$

At training time, the objective is standard noise-prediction MSE:

$$\mathcal{L}(\theta) = \mathbb{E}_{x_0, \epsilon, t} \left[ \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|_2^2 \right].$$

## 4 Network Architecture

The current model is a residual U-Net style denoiser with sinusoidal time embeddings:

- Stem: Conv 1 → 64
- Down path: residual blocks at 64 and 128 channels with strided downsampling (28 → 14 → 7)
- Bottleneck: two residual blocks at 256 channels
- Up path: transposed convolutions + skip concatenation + residual blocks (7 → 14 → 28)
- Output: GroupNorm + Conv 64 → 1
- Optimizer: AdamW, gradient clipping (max norm 1.0)

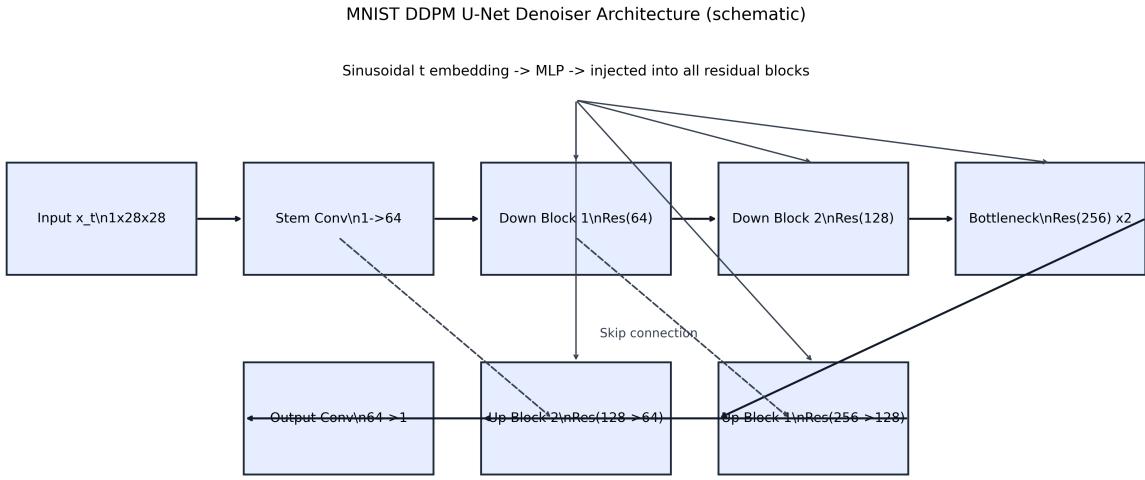


Figure 1: U-Net denoiser schematic generated by the training pipeline.

## 5 Latest Long Run Configuration

The latest long run (job 8109577) used:

- Epochs: 80
- Diffusion steps ( $T$ ): 300
- Batch size: 128
- Slurm resources: 1 H100 GPU, 14 CPUs, partition `ai`, walltime 4h

From `metrics.json` (run `long_8109577`):

- Total steps: 37,520
- Final epoch mean loss: 0.0362516
- Best epoch mean loss: 0.0360608
- Device: CUDA

## 6 Results and Artifacts

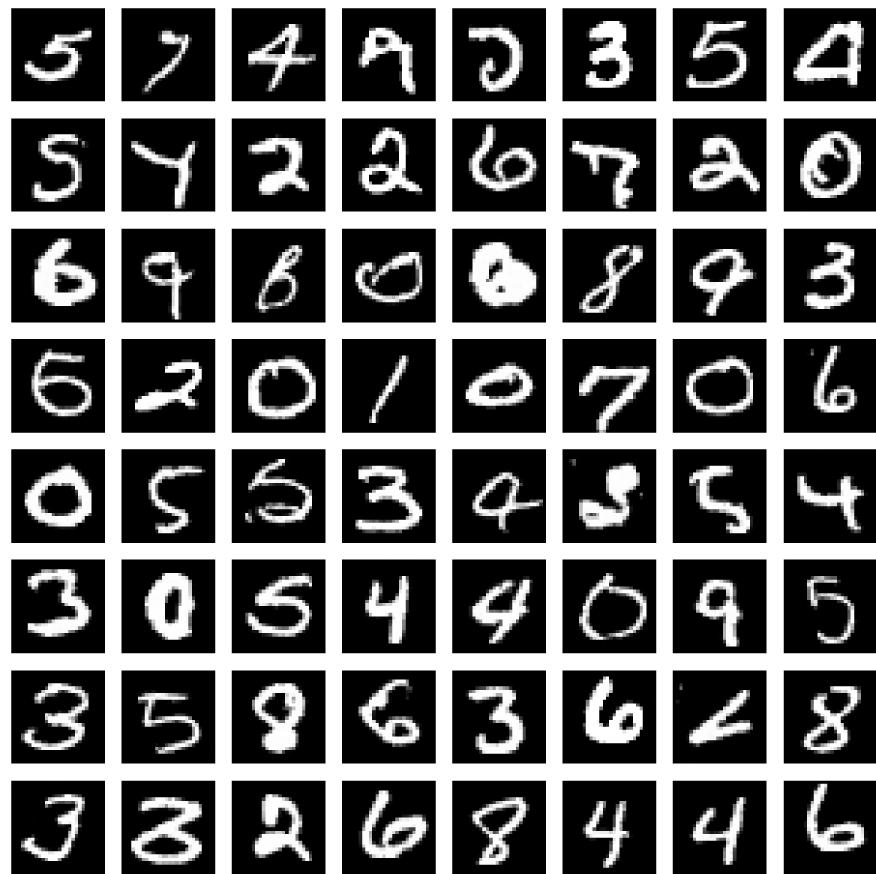


Figure 2: Generated MNIST samples from the latest completed run.

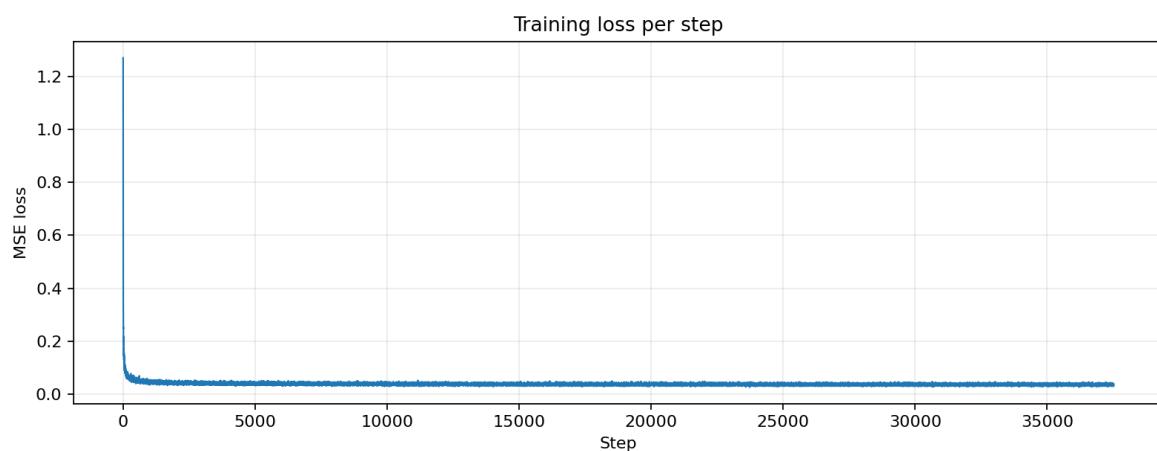


Figure 3: Per-step training loss trajectory.

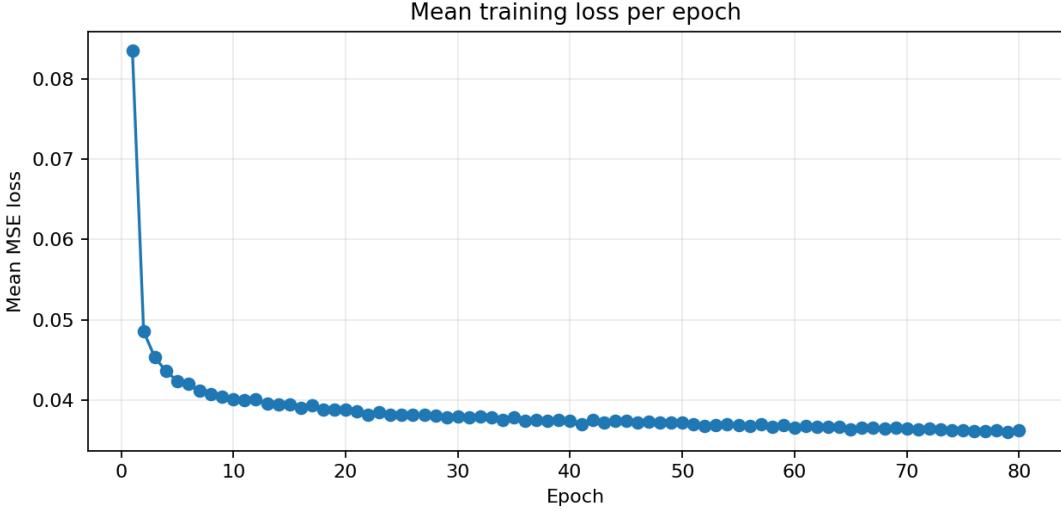


Figure 4: Per-epoch mean loss trajectory.

## 7 Posterior Sampling with Partial Observations

We added a post-training posterior sampling routine to `mnist_diffusion.py` that performs likelihood-guided reverse diffusion under partial pixel observations. Let  $y$  denote observed pixels,  $m \in \{0, 1\}^{H \times W}$  the mask, and  $\hat{x}_0(x_t, t)$  the model-implied clean image estimate.

### Mathematical Formulation

At each reverse step, the baseline DDPM mean is:

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right).$$

Posterior guidance is applied through a masked Gaussian likelihood:

$$p(y | x_0) \propto \exp \left( -\frac{\|m \odot (y - x_0)\|_2^2}{2\sigma_t^2} \right), \quad \sigma_t^2 = \sigma_y^2 + c(1 - \bar{\alpha}_t).$$

This yields the guidance gradient (used in code):

$$g_t \approx \frac{m \odot (y - \hat{x}_0)}{\sigma_t^2 \sqrt{\bar{\alpha}_t}}.$$

The reverse mean is then updated as:

$$\tilde{\mu}_t = \mu_\theta(x_t, t) + \lambda_t \beta_t g_t,$$

with timestep-annealed guidance strength

$$\lambda_t = \lambda_{\max} \left( \lambda_{\min} + (1 - \lambda_{\min}) \left( 1 - \frac{t}{T-1} \right)^p \right).$$

We additionally enforce projection-based data consistency:

$$\hat{x}_0 \leftarrow m \odot y + (1 - m) \odot \hat{x}_0,$$

and for  $t > 0$ :

$$x_{t-1} \leftarrow m \odot \left( \sqrt{\bar{\alpha}_{t-1}} y + \sqrt{1 - \bar{\alpha}_{t-1}} \xi \right) + (1 - m) \odot x_{t-1}.$$

The implemented strategy therefore combines:

- Noise-aware likelihood variance:  $\sigma_t^2 = \sigma_y^2 + c(1 - \bar{\alpha}_t)$
- Annealed guidance schedule: weak guidance at high noise, stronger guidance near final denoise steps
- Data-consistency projection on observed pixels (in both  $\hat{x}_0$  and sampled states)

### Experiment A: 70% Pixels Revealed

- Observed fraction: 70% of pixels (30% occluded, random mask)
- Target digit class: 7
- Guidance scale: 1.5
- Guidance min fraction / power: 0.25 / 1.5
- Likelihood noise scale: 0.1
- Noise-aware coefficient: 0.05

Validated Gautschi posterior job:

- Job ID: 8112324
- State: COMPLETED, exit code 0:0
- Elapsed: 00:00:27
- Run tag: posterior\_70pct\_dpsplus\_8112324

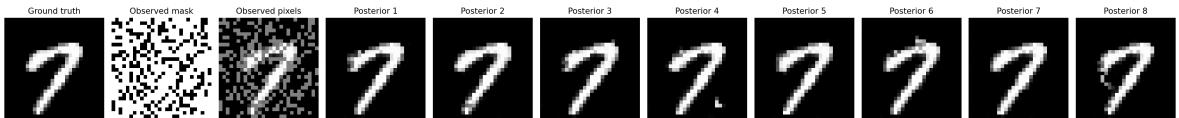


Figure 5: Improved DPS conditioning setup: ground truth, observed mask (70% observed), observed pixels, and posterior draws.



Figure 6: Posterior samples with noise-aware likelihood, guidance annealing, and projection-based data consistency (70% revealed).

## Experiment B: 20% Pixels Revealed

- Observed fraction: 20% of pixels (80% occluded, random mask)
- Target digit class: 7
- Guidance scale: 1.5
- Guidance min fraction / power: 0.25 / 1.5
- Likelihood noise scale: 0.1
- Noise-aware coefficient: 0.05

Validated Gautschi posterior job:

- Job ID: 8112347
- State: COMPLETED, exit code 0:0
- Elapsed: 00:00:28
- Run tag: posterior\_20pct\_dpsplus\_8112347

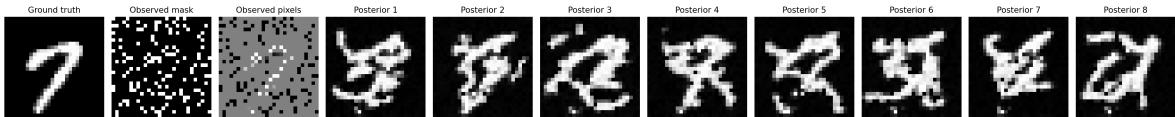


Figure 7: Improved DPS conditioning setup: ground truth, observed mask (20% observed), observed pixels, and posterior draws.



Figure 8: Posterior samples with noise-aware likelihood, guidance annealing, and projection-based data consistency (20% revealed).

## 8 Live Dashboard Operation

The workflow now maintains a root live dashboard at:

- `outputs/dashboard.html`

This dashboard follows the latest run via:

- `outputs/LATEST_RUN.txt`
- `outputs/current` symlink

and supports:

- periodic refresh during active runs,
- automatic stop of refresh when status reaches `completed`,
- image zoom controls for plots/schematics.

## 9 Reproducibility Notes

Primary scripts tracked in Git:

- `mnist_diffusion.py`
- `submit_mnist_diffusion_long.slurm`
- `serve_dashboard.sh`

### Persistent SSH Connection Reuse

To avoid repeated password + 2FA prompts for each `ssh/scp` call, use SSH multiplexing with a persistent control socket:

```
ssh -M -S /tmp/gautschi_mux.sock -o ControlPersist=8h -N \
rmaulik@gautschi.rcac.purdue.edu
```

Then reuse that authenticated channel:

```
ssh -S /tmp/gautschi_mux.sock rmaulik@gautschi.rcac.purdue.edu
scp -o ControlMaster=auto -o ControlPath=/tmp/gautschi_mux.sock \
<src> <dst>
```

Check and close the master session:

```
ssh -S /tmp/gautschi_mux.sock -O check rmaulik@gautschi.rcac.purdue.edu
ssh -S /tmp/gautschi_mux.sock -O exit rmaulik@gautschi.rcac.purdue.edu
```

To regenerate this report PDF locally:

```
cd PurdueHPC_Codex/report
latexmk -pdf -interaction=nonstopmode -halt-on-error \
-output-directory=../output/pdf gautschi_diffusion_report.tex
```