Technical Report: MNIST Diffusion Experiments on Purdue Gautschi

PurdueHPC_Codex Workflow

March 1, 2026

1 Scope and Environment

This report summarizes end-to-end execution of diffusion model experiments on Purdue RCAC Gautschi, including environment setup, Slurm execution, model details, and resulting artifacts. All commands and scripts were executed under:

- Scratch workspace: `/scratch/gautschi/rmaulik/codex_test`
- Source repository: `PurdueHPC_Codex`
- Account: `rmaulik`
- Partition: `ai`

No credentials or authentication secrets are stored in this report.

2 Allocation and Job Accounting

Current allocation snapshot from `slist` on Gautschi:

- AI partition GPU-hour balance: **43,798.5**
- CPU partition balance: **0**

Completed follow-up experiment jobs (EDM formulation):

Job ID	Name	State	Elapsed (s)	GPUs
8213207	mnist_edm	COMPLETED	28	1
8213218	mnist_edm20	COMPLETED	30	1

Estimated direct GPU-hours from these two runs:

$$\text{GPU-hours} = \frac{28 + 30}{3600} \times 1 = 0.016 \text{ GPU-hours (approx.)}$$

3 Diffusion Model Formulation (EDM)

The training script now uses an Elucidated Diffusion Model (EDM) style preconditioned denoiser over MNIST ($x_0 \in [-1, 1]^{1 \times 28 \times 28}$), where the model approximates a score-consistent denoising function across continuous noise levels σ.

3.1 EDM Preconditioning and Objective

Given $x = x_0 + \sigma\epsilon$ with $\epsilon \sim \mathcal{N}(0, I)$ and $\sigma \sim \log\mathcal{N}(p_{\text{mean}}, p_{\text{std}}^2)$, the network is used in preconditioned form:

$$\hat{x}_0 = c_{\text{skip}}(\sigma)x + c_{\text{out}}(\sigma)F_\theta(c_{\text{in}}(\sigma)x, c_{\text{noise}}(\sigma)),$$

where

$$c_{\text{in}} = \frac{1}{\sqrt{\sigma^2 + \sigma_{\text{data}}^2}}, \quad c_{\text{skip}} = \frac{\sigma_{\text{data}}^2}{\sigma^2 + \sigma_{\text{data}}^2}, \quad c_{\text{out}} = \frac{\sigma\sigma_{\text{data}}}{\sqrt{\sigma^2 + \sigma_{\text{data}}^2}}.$$

Training minimizes weighted denoising MSE:

$$\mathcal{L}(\theta) = \mathbb{E}_{x_0, \sigma, \epsilon}\left[w(\sigma)\|\hat{x}_0 - x_0\|_2^2\right], \quad w(\sigma) = \frac{\sigma^2 + \sigma_{\text{data}}^2}{(\sigma\sigma_{\text{data}})^2}.$$

3.2 EDM Sampling

Sampling follows a Karras σ-schedule from σ_{max} to σ_{min} (then 0) and integrates the probability flow ODE in σ-space:

$$\frac{dx}{d\sigma} = \frac{x - \hat{x}_0(x, \sigma)}{\sigma},$$

with Euler step plus second-order (Heun) correction.

4 Network Architecture

The current model is a residual U-Net style denoiser with sinusoidal time embeddings:

- Stem: Conv $1 \rightarrow 64$

- Down path: residual blocks at 64 and 128 channels with strided downsampling ($28 \rightarrow 14 \rightarrow 7$)

- Bottleneck: two residual blocks at 256 channels

- Up path: transposed convolutions + skip concatenation + residual blocks ($7 \rightarrow 14 \rightarrow 28$)

- Output: GroupNorm + Conv $64 \rightarrow 1$

- Optimizer: AdamW, gradient clipping (max norm 1.0)

MNIST EDM U-Net Denoiser Architecture (schematic)

Sinusoidal t embedding -> MLP -> injected into all residual blocks

Input x_sigma\n1x28x28 → Stem Conv\n1->64 → Down Block 1\nRes(64) → Down Block 2\nRes(128) → Bottleneck\nRes(256) x2

Skip connection

Output Conv\n64->1 ← Up Block 2\nRes(128->64) ← Up Block 1\nRes(256->128)

Figure 1: U-Net denoiser schematic generated by the EDM training pipeline.

5 Latest EDM Run Configuration

The latest EDM follow-up run (job 8213207) used:

- Epochs: 1

- Karras sampling steps: 80

- Batch size: 128

- EDM parameters: $\sigma_{\min} = 0.002$, $\sigma_{\max} = 80.0$, $\rho = 7.0$, $\sigma_{\text{data}} = 0.5$

- Sigma training distribution: $p_{\text{mean}} = -1.2$, $p_{\text{std}} = 1.2$

- Slurm resources: 1 GPU, 14 CPUs, partition ai

From metrics.json (run edm_dps_70pct_8213207):

- Total steps: 469

- Final epoch mean loss: 0.273500

- Best epoch mean loss: 0.273500

- Device: CUDA

6   Results and Artifacts

Figure 2: Generated MNIST samples from the EDM run (8213207).


Figure 3: Per-step training loss trajectory for EDM run 8213207.

Figure 4: Per-epoch mean loss trajectory for EDM run 8213207.

7 Posterior Sampling with Partial Observations

We perform diffusion posterior sampling in EDM sigma-space. Let y denote observed pixels, $m \in \{0,1\}^{H \times W}$ the mask, and $\hat{x}_0(x, \sigma)$ the EDM denoised estimate.

Mathematical Formulation

Posterior guidance is applied through a masked Gaussian likelihood:

$$p(y \mid x_0) \propto \exp\left(-\frac{\|m \odot (y - x_0)\|_2^2}{2\sigma_t^2}\right), \quad \sigma_t^2 = \sigma_y^2 + c\sigma^2.$$

This yields the guidance gradient (used in code):

$$g_\sigma \approx \frac{m \odot (y - \hat{x}_0)}{\sigma_t^2}.$$

In EDM ODE form, we use:

$$\frac{dx}{d\sigma} \approx \frac{x - (\hat{x}_0 + \lambda_\sigma \sigma g_\sigma)}{\sigma},$$

with timestep-annealed guidance strength

$$\lambda_\sigma = \lambda_{\max}\left(\lambda_{\min} + (1 - \lambda_{\min})\left(\frac{i}{N - 1}\right)^p\right).$$

We additionally enforce projection-based data consistency:

$$\hat{x}_0 \leftarrow m \odot y + (1 - m) \odot \hat{x}_0,$$

and during intermediate sigma steps:

$$x \leftarrow m \odot (y + \sigma\xi) + (1 - m) \odot x.$$

The implemented strategy therefore combines:

- Noise-aware likelihood variance: $\sigma_t^2 = \sigma_y^2 + c\sigma^2$

- Annealed guidance schedule: weak guidance at high noise, stronger guidance near final denoise steps

- Data-consistency projection on observed pixels (in both $\hat{x}_0$ and sampled states)

Experiment A: 70% Pixels Revealed

- Observed fraction: 70% of pixels (30% occluded, random mask)

- Target digit class: 7

- Guidance scale: 1.5

- Guidance min fraction / power: 0.25 / 1.5

- Likelihood noise scale: 0.1

- Noise-aware coefficient: 0.05

Validated Gautschi posterior job:

- Job ID: `8213207`

- State: `COMPLETED`, exit code `0:0`

- Elapsed: `00:00:28`

- Run tag: `edm_dps_70pct_8213207`

Figure 5: EDM posterior conditioning setup: ground truth, observed mask (70% observed), observed pixels, and posterior draws.

Figure 6: EDM posterior samples with noise-aware likelihood, guidance annealing, and projection-based data consistency (70% revealed).

Experiment B: 20% Pixels Revealed

- Observed fraction: 20% of pixels (80% occluded, random mask)

- Target digit class: 7

- Guidance scale: 1.5

- Guidance min fraction / power: 0.25 / 1.5

- Likelihood noise scale: 0.1

- Noise-aware coefficient: 0.05

Validated Gautschi posterior job:

- Job ID: `8213218`

- State: `COMPLETED`, exit code `0:0`

- Elapsed: `00:00:30`

- Run tag: `edm_dps_20pct_8213218`

Figure 7: EDM posterior conditioning setup: ground truth, observed mask (20% observed), observed pixels, and posterior draws.

Figure 8: EDM posterior samples with noise-aware likelihood, guidance annealing, and projection-based data consistency (20% revealed).

8 Intrinsic Dimensionality of Posterior Samples

To analyze posterior sample complexity, we generated a large posterior ensemble for each setting and formed a snapshot matrix:

$$X = [x^{(1)}, x^{(2)}, \ldots, x^{(N)}] \in \mathbb{R}^{784 \times N},$$

where each column is a flattened 28×28 posterior sample and $N = 1024 > 784$. We computed

$$X = U \Sigma V^\top,$$

and inspected the singular value spectrum $\{\sigma_i\}$ and cumulative energy.

SVD jobs on Gautschi:

- 70% observed: job `8213341`, run `edm_svd_70pct_8213341`

- 20% observed: job `8213342`, run `edm_svd_20pct_8213342`

Observed fraction	Pixels	Snapshots	k_{95}	k_{99}	Participation ratio
70%	784	1024	1	3	1.023
20%	784	1024	45	152	1.718

Table 2: Singular-spectrum summary from posterior snapshot matrices.
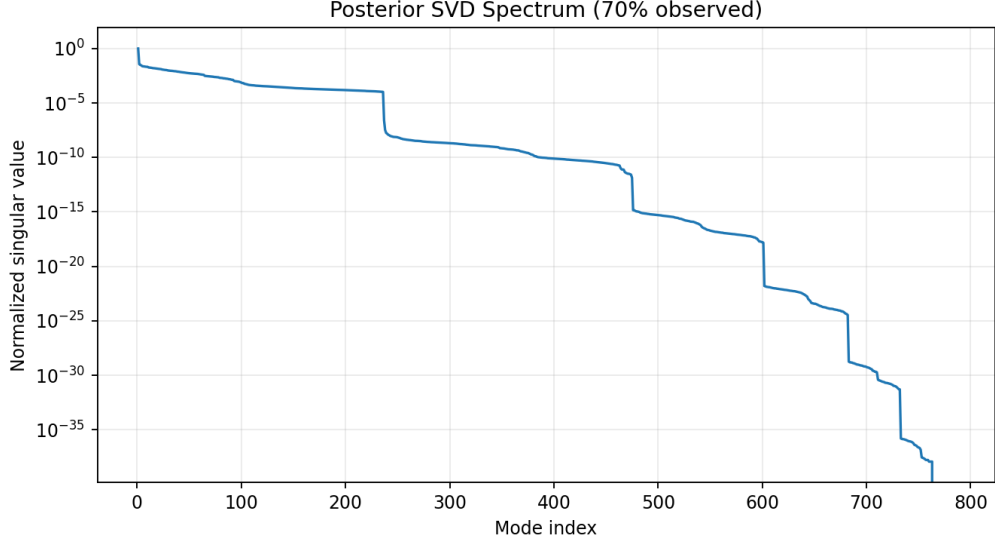
Figure 9: Normalized singular value spectrum for EDM posterior samples with 70% observed pixels.
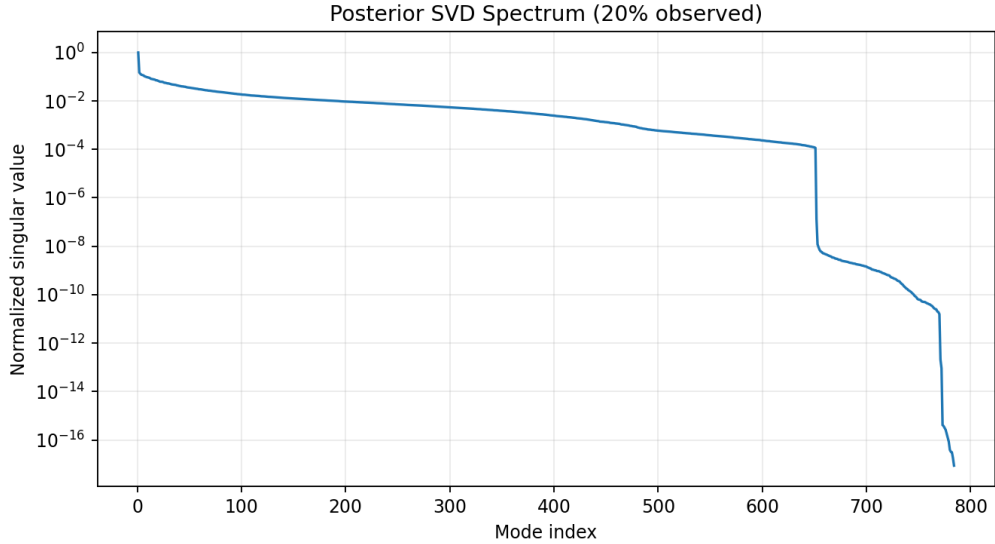

Figure 10: Normalized singular value spectrum for EDM posterior samples with 20% observed pixels.

Additional Subsection: Score-Snapshot SVD and Spectrum Overlay

We repeated the intrinsic dimensionality analysis by constructing snapshots from the *estimated score field* instead of the samples. For each posterior sample $x^{(j)}$, we evaluated the approximated score at fixed $\sigma_{\text{eval}} = 0.1$:

$$s^{(j)} \approx \frac{\hat{x}_0(x^{(j)}, \sigma_{\text{eval}}) - x^{(j)}}{\sigma_{\text{eval}}^2},$$

then formed
$$S = [s^{(1)}, s^{(2)}, \ldots, s^{(N)}] \in \mathbb{R}^{784 \times N},$$

with $N = 1024 > 784$, and computed its SVD.

Score-snapshot jobs on Gautschi:

- 70% observed: job `8213437`, run `edm_score_svd_70pct_8213437`

- 20% observed: job `8213438`, run `edm_score_svd_20pct_8213438`

Observed fraction	Snapshot type	k_{95}	k_{99}	Participation ratio
70%	Sample (previous)	1	2	1.023
70%	Score (current)	63	137	2.566
20%	Sample (previous)	46	154	1.735
20%	Score (current)	293	464	8.786

Table 3: Comparison of singular-spectrum intrinsic dimensionality metrics for sample- vs score-based snapshots.


Figure 11: Overlay of normalized singular value spectra (70% observed): previous sample-snapshot experiment vs current score-snapshot experiment.

Figure 12: Overlay of normalized singular value spectra (20% observed): previous sample-snapshot experiment vs current score-snapshot experiment.

9 Live Dashboard Operation

The workflow now maintains a root live dashboard at:

- `outputs/dashboard.html`

This dashboard follows the latest run via:

- `outputs/LATEST_RUN.txt`
- `outputs/current` symlink

and supports:

- periodic refresh during active runs,
- automatic stop of refresh when status reaches `completed`,
- image zoom controls for plots/schematics.

10 Reproducibility Notes

Primary scripts used for these experiments:

- `mnist_diffusion.py`
- `posterior_svd_analysis.py`
- `posterior_score_svd_analysis.py`
- `submit_mnist_edm_posterior.slurm`
- `submit_mnist_edm_posterior_20pct.slurm`
- `serve_dashboard.sh`

Persistent SSH Connection Reuse

To avoid repeated password + 2FA prompts for each `ssh`/`scp` call, use SSH multiplexing with a persistent control socket:

```
ssh -M -S /tmp/gautschi_mux.sock -o ControlPersist=8h -N \
  rmaulik@gautschi.rcac.purdue.edu
```

Then reuse that authenticated channel:

```
ssh -S /tmp/gautschi_mux.sock rmaulik@gautschi.rcac.purdue.edu
scp -o ControlMaster=auto -o ControlPath=/tmp/gautschi_mux.sock \
  <src> <dst>
```

Check and close the master session:

```
ssh -S /tmp/gautschi_mux.sock -O check rmaulik@gautschi.rcac.purdue.edu
ssh -S /tmp/gautschi_mux.sock -O exit  rmaulik@gautschi.rcac.purdue.edu
```

To regenerate this report PDF locally:

```
cd PurdueHPC_Codex/report
latexmk -pdf -interaction=nonstopmode -halt-on-error \
  -output-directory=../output/pdf gautschi_diffusion_report.tex
```