

Technical Report: MNIST Diffusion Experiments on Purdue Gautschi

PurdueHPC_Codex Workflow

March 1, 2026

1 Scope and Environment

This report summarizes end-to-end execution of diffusion model experiments on Purdue RCAC Gautschi, including environment setup, Slurm execution, model details, and resulting artifacts. All commands and scripts were executed under:

- Scratch workspace: `/scratch/gautschi/rmaulik/codex_test`
- Source repository: `PurdueHPC_Codex`
- Account: `rmaulik`
- Partition: `ai`

No credentials or authentication secrets are stored in this report.

2 Allocation and Job Accounting

Current allocation snapshot from `slist` on Gautschi:

- AI partition GPU-hour balance: **43,798.5**
- CPU partition balance: **0**

Completed follow-up experiment jobs (EDM formulation):

Job ID	Name	State	Elapsed (s)	GPUs
8213207	mnist_edm	COMPLETED	28	1
8213218	mnist_edm20	COMPLETED	30	1

Estimated direct GPU-hours from these two runs:

$$\text{GPU-hours} = \frac{28 + 30}{3600} \times 1 = 0.016 \text{ GPU-hours (approx.)}$$

3 Diffusion Model Formulation (EDM)

The training script now uses an Elucidated Diffusion Model (EDM) style preconditioned denoiser over MNIST ($x_0 \in [-1, 1]^{1 \times 28 \times 28}$), where the model approximates a score-consistent denoising function across continuous noise levels σ .

3.1 EDM Preconditioning and Objective

Given $x = x_0 + \sigma\epsilon$ with $\epsilon \sim \mathcal{N}(0, I)$ and $\sigma \sim \log \mathcal{N}(p_{\text{mean}}, p_{\text{std}}^2)$, the network is used in preconditioned form:

$$\hat{x}_0 = c_{\text{skip}}(\sigma)x + c_{\text{out}}(\sigma)F_\theta(c_{\text{in}}(\sigma)x, c_{\text{noise}}(\sigma)),$$

where

$$c_{\text{in}} = \frac{1}{\sqrt{\sigma^2 + \sigma_{\text{data}}^2}}, \quad c_{\text{skip}} = \frac{\sigma_{\text{data}}^2}{\sigma^2 + \sigma_{\text{data}}^2}, \quad c_{\text{out}} = \frac{\sigma\sigma_{\text{data}}}{\sqrt{\sigma^2 + \sigma_{\text{data}}^2}}.$$

Training minimizes weighted denoising MSE:

$$\mathcal{L}(\theta) = \mathbb{E}_{x_0, \sigma, \epsilon} [w(\sigma)\|\hat{x}_0 - x_0\|_2^2], \quad w(\sigma) = \frac{\sigma^2 + \sigma_{\text{data}}^2}{(\sigma\sigma_{\text{data}})^2}.$$

3.2 EDM Sampling

Sampling follows a Karras σ -schedule from σ_{max} to σ_{min} (then 0) and integrates the probability flow ODE in σ -space:

$$\frac{dx}{d\sigma} = \frac{x - \hat{x}_0(x, \sigma)}{\sigma},$$

with Euler step plus second-order (Heun) correction.

4 Network Architecture

The current model is a residual U-Net style denoiser with sinusoidal time embeddings:

- Stem: Conv 1 → 64
- Down path: residual blocks at 64 and 128 channels with strided downsampling (28 → 14 → 7)
- Bottleneck: two residual blocks at 256 channels
- Up path: transposed convolutions + skip concatenation + residual blocks (7 → 14 → 28)
- Output: GroupNorm + Conv 64 → 1
- Optimizer: AdamW, gradient clipping (max norm 1.0)

MNIST EDM U-Net Denoiser Architecture (schematic)

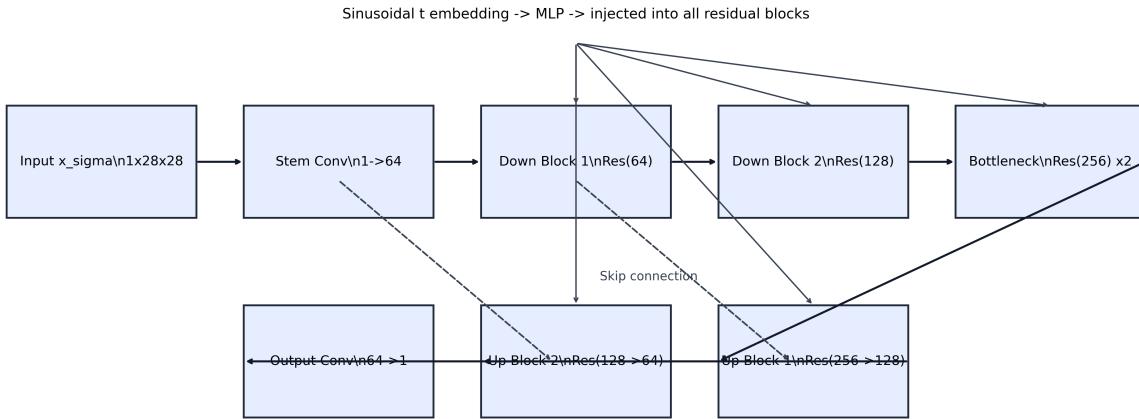


Figure 1: U-Net denoiser schematic generated by the EDM training pipeline.

5 Latest EDM Run Configuration

The latest EDM follow-up run (job 8213207) used:

- Epochs: 1
- Karras sampling steps: 80
- Batch size: 128
- EDM parameters: $\sigma_{\min} = 0.002$, $\sigma_{\max} = 80.0$, $\rho = 7.0$, $\sigma_{\text{data}} = 0.5$
- Sigma training distribution: $p_{\text{mean}} = -1.2$, $p_{\text{std}} = 1.2$
- Slurm resources: 1 GPU, 14 CPUs, partition **ai**

From `metrics.json` (run `edm_dps_70pct_8213207`):

- Total steps: 469
- Final epoch mean loss: 0.273500
- Best epoch mean loss: 0.273500
- Device: CUDA

6 Results and Artifacts

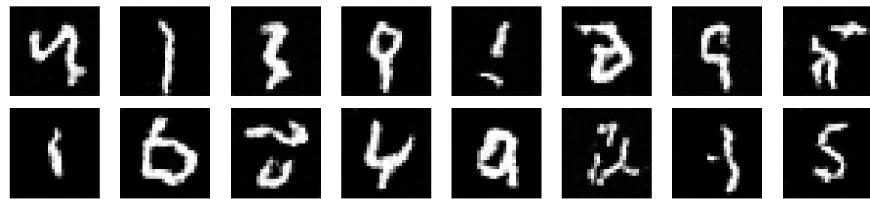


Figure 2: Generated MNIST samples from the EDM run (8213207).

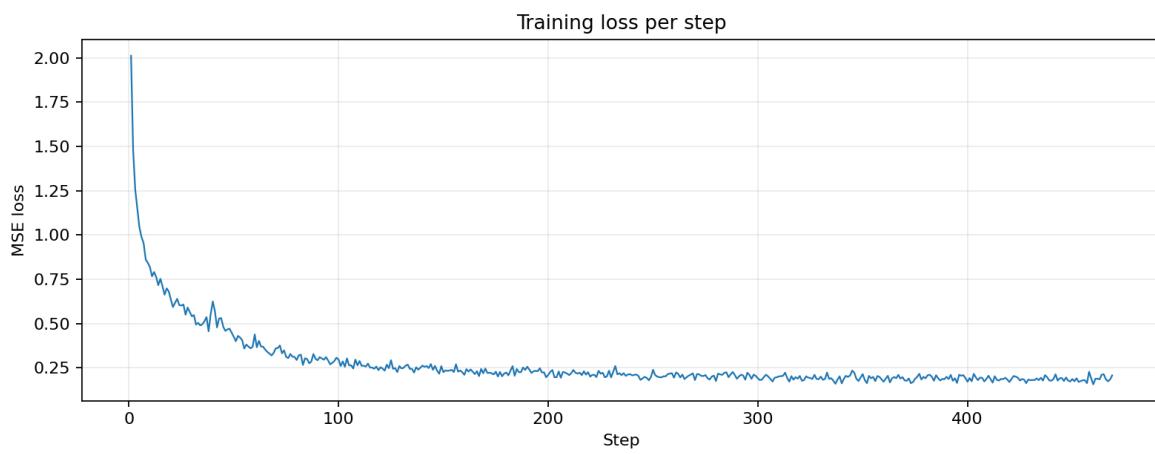


Figure 3: Per-step training loss trajectory for EDM run 8213207.

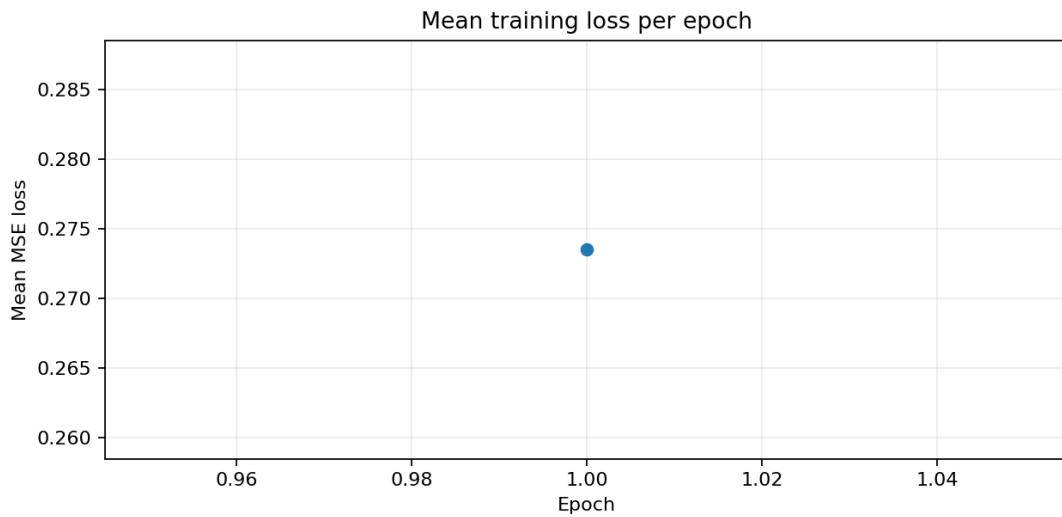


Figure 4: Per-epoch mean loss trajectory for EDM run 8213207.

7 Posterior Sampling with Partial Observations

We perform diffusion posterior sampling in EDM sigma-space. Let y denote observed pixels, $m \in \{0, 1\}^{H \times W}$ the mask, and $\hat{x}_0(x, \sigma)$ the EDM denoised estimate.

Mathematical Formulation

Posterior guidance is applied through a masked Gaussian likelihood:

$$p(y | x_0) \propto \exp\left(-\frac{\|m \odot (y - x_0)\|_2^2}{2\sigma_t^2}\right), \quad \sigma_t^2 = \sigma_y^2 + c\sigma^2.$$

This yields the guidance gradient (used in code):

$$g_\sigma \approx \frac{m \odot (y - \hat{x}_0)}{\sigma_t^2}.$$

In EDM ODE form, we use:

$$\frac{dx}{d\sigma} \approx \frac{x - (\hat{x}_0 + \lambda_\sigma \sigma g_\sigma)}{\sigma},$$

with timestep-annealed guidance strength

$$\lambda_\sigma = \lambda_{\max} \left(\lambda_{\min} + (1 - \lambda_{\min}) \left(\frac{i}{N-1} \right)^p \right).$$

We additionally enforce projection-based data consistency:

$$\hat{x}_0 \leftarrow m \odot y + (1 - m) \odot \hat{x}_0,$$

and during intermediate sigma steps:

$$x \leftarrow m \odot (y + \sigma \xi) + (1 - m) \odot x.$$

The implemented strategy therefore combines:

- Noise-aware likelihood variance: $\sigma_t^2 = \sigma_y^2 + c\sigma^2$
- Annealed guidance schedule: weak guidance at high noise, stronger guidance near final denoise steps
- Data-consistency projection on observed pixels (in both \hat{x}_0 and sampled states)

Experiment A: 70% Pixels Revealed

- Observed fraction: 70% of pixels (30% occluded, random mask)
- Target digit class: 7
- Guidance scale: 1.5
- Guidance min fraction / power: 0.25 / 1.5
- Likelihood noise scale: 0.1

- Noise-aware coefficient: 0.05

Validated Gautschi posterior job:

- Job ID: 8213207
- State: COMPLETED, exit code 0:0
- Elapsed: 00:00:28
- Run tag: `edm_dps_70pct_8213207`

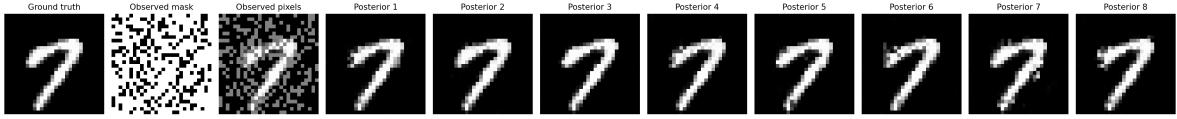


Figure 5: EDM posterior conditioning setup: ground truth, observed mask (70% observed), observed pixels, and posterior draws.



Figure 6: EDM posterior samples with noise-aware likelihood, guidance annealing, and projection-based data consistency (70% revealed).

Experiment B: 20% Pixels Revealed

- Observed fraction: 20% of pixels (80% occluded, random mask)
- Target digit class: 7
- Guidance scale: 1.5
- Guidance min fraction / power: 0.25 / 1.5
- Likelihood noise scale: 0.1
- Noise-aware coefficient: 0.05

Validated Gautschi posterior job:

- Job ID: 8213218
- State: COMPLETED, exit code 0:0
- Elapsed: 00:00:30
- Run tag: `edm_dps_20pct_8213218`

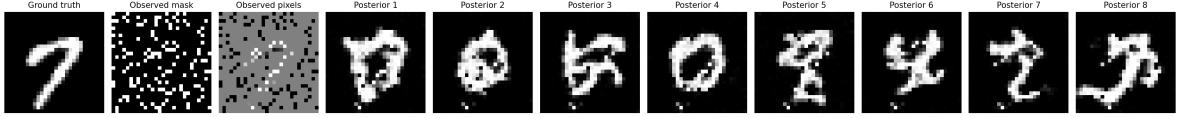


Figure 7: EDM posterior conditioning setup: ground truth, observed mask (20% observed), observed pixels, and posterior draws.



Figure 8: EDM posterior samples with noise-aware likelihood, guidance annealing, and projection-based data consistency (20% revealed).

8 Live Dashboard Operation

The workflow now maintains a root live dashboard at:

- `outputs/dashboard.html`

This dashboard follows the latest run via:

- `outputs/LATEST_RUN.txt`
- `outputs/current` symlink

and supports:

- periodic refresh during active runs,
- automatic stop of refresh when status reaches `completed`,
- image zoom controls for plots/schematics.

9 Reproducibility Notes

Primary scripts used for these experiments:

- `mnist_diffusion.py`
- `submit_mnist_edm_posterior.slurm`
- `submit_mnist_edm_posterior_20pct.slurm`
- `serve_dashboard.sh`

Persistent SSH Connection Reuse

To avoid repeated password + 2FA prompts for each `ssh/scp` call, use SSH multiplexing with a persistent control socket:

```
ssh -M -S /tmp/gautschi_mux.sock -o ControlPersist=8h -N \
rmaulik@gautschi.rcac.purdue.edu
```

Then reuse that authenticated channel:

```
ssh -S /tmp/gautschi_mux.sock rmaulik@gautschi.rcac.purdue.edu
scp -o ControlMaster=auto -o ControlPath=/tmp/gautschi_mux.sock \
<src> <dst>
```

Check and close the master session:

```
ssh -S /tmp/gautschi_mux.sock -O check rmaulik@gautschi.rcac.purdue.edu
ssh -S /tmp/gautschi_mux.sock -O exit rmaulik@gautschi.rcac.purdue.edu
```

To regenerate this report PDF locally:

```
cd PurdueHPC_Codex/report
latexmk -pdf -interaction=nonstopmode -halt-on-error \
-output-directory=../output/pdf gautschi_diffusion_report.tex
```